

Aaron Romero

CSE 160

To complete the routing project, I started by refactoring the flooding module to be able to send a packet with an input protocol. Then I just made a switch case statement over the input protocol, and for this project the only one that matters is `PROTOCOL_LINKSTATE`. All it does is use the input protocol as the protocol header for the packet. Also, I made it so that `AM_BROADCAST_ADDR` can be used to make all nodes receive packets, read them, and then pass them on via flooding. Once this was finished, I had to use the `NeighborDiscovery.neighbors` command to help the Linkstate module know what neighbors it has. I compressed this information into one bit per neighbor, so if there are 40 nodes then 5 bytes will be stored with each bit representing whether Node X is a neighbor. Then I just used these 5 bytes as the message for the flood packet, and set each node to start sending out neighbor information packets 100 seconds after `start()` is called using a timer. Once the packet is received, it goes into `Receive.receive` in `node.nc`, which reads the protocol in a switch case. In the case of `protocol_linkstate`, we call `Flooding.handleFlood`, which will see that it has a linkstate packet, and send it to the linkstate module with `linkstate.receiveUpdate`, which will then update each node's 2D array of neighbor values. This array is length  $N * \text{ceil}(N / 8)$ , where N is the number of nodes. Note that  $\text{ceil}(N/8)$  is the number of bytes required to store neighbor values for a singular node. Once all neighbor values are computed, we call `dijkstra's` algorithm and create a routing table which simply dictates which neighbor packets are sent to on the next hop, and this logic is managed in `LinkState.handleRoutingPacket`. I also use a second protocol called `DIRECTROUTE_PROTOCOL` for packets which will need to be directly routed rather than flooded.